

Process Management: Processes

Lecturer: Kor Sokchea

Royal University of Phnom Penh

Faculty of Engineering

kor.sokchea@gmail.com



- **Textbook:** *Operating System Concepts*, 9th Edition, by Silberschatz, Galvin, and Gagne.
- Class
 - IT Engineering: Room **T002**
 - Monday Session: 13:00 - 14:30
 - Tuesday Session: 16:00 - 17:30
 - TEED: Room **T104**
 - Monday Session: 16:00 - 17:30
 - Tuesday Session: 13:00 - 14:30
- Exams
 - Final Exam: 60%
 - Homework: 10%
 - Assignment: 20%
 - Attendance: 10%

Course Topics

- **Threads & Processes**
- **Concurrency & Synchronization**
- **Scheduling**
- **Deadlocks**
- **Main Memory**
- **Virtual Memory**
- **I/O systems**
- **Disk**
- **File System**
- **Protection & Security**

Process Management

- **Process**
- **Threads**
- **Process Synchronization**
- **CPU Scheduling**

What is a Process?

- A process is an instance of a program running
 - Program = static file (image)
 - Process = executing program + execution state
- A Process is the basic unit of execution in an operating system
 - Each process has a number, its process identifier (**pid**)
- Different processes may run different instances of the same program
 - Ex: my gcc and your gcc process both run c compiler
- At a minimum, process execution requires following resources:
 - Memory to contain the program code and data
 - A set of CPU registers to support execution

What is a Process?

- Modern OS run multiple processes concurrently
- Examples: (can all run simultaneously):
 - gcc hello_world.c – compiler running on **hello_world** file
 - gcc read_A.c – compiler running on **read_A** file
 - emacs – text editor
 - firefox – web browser

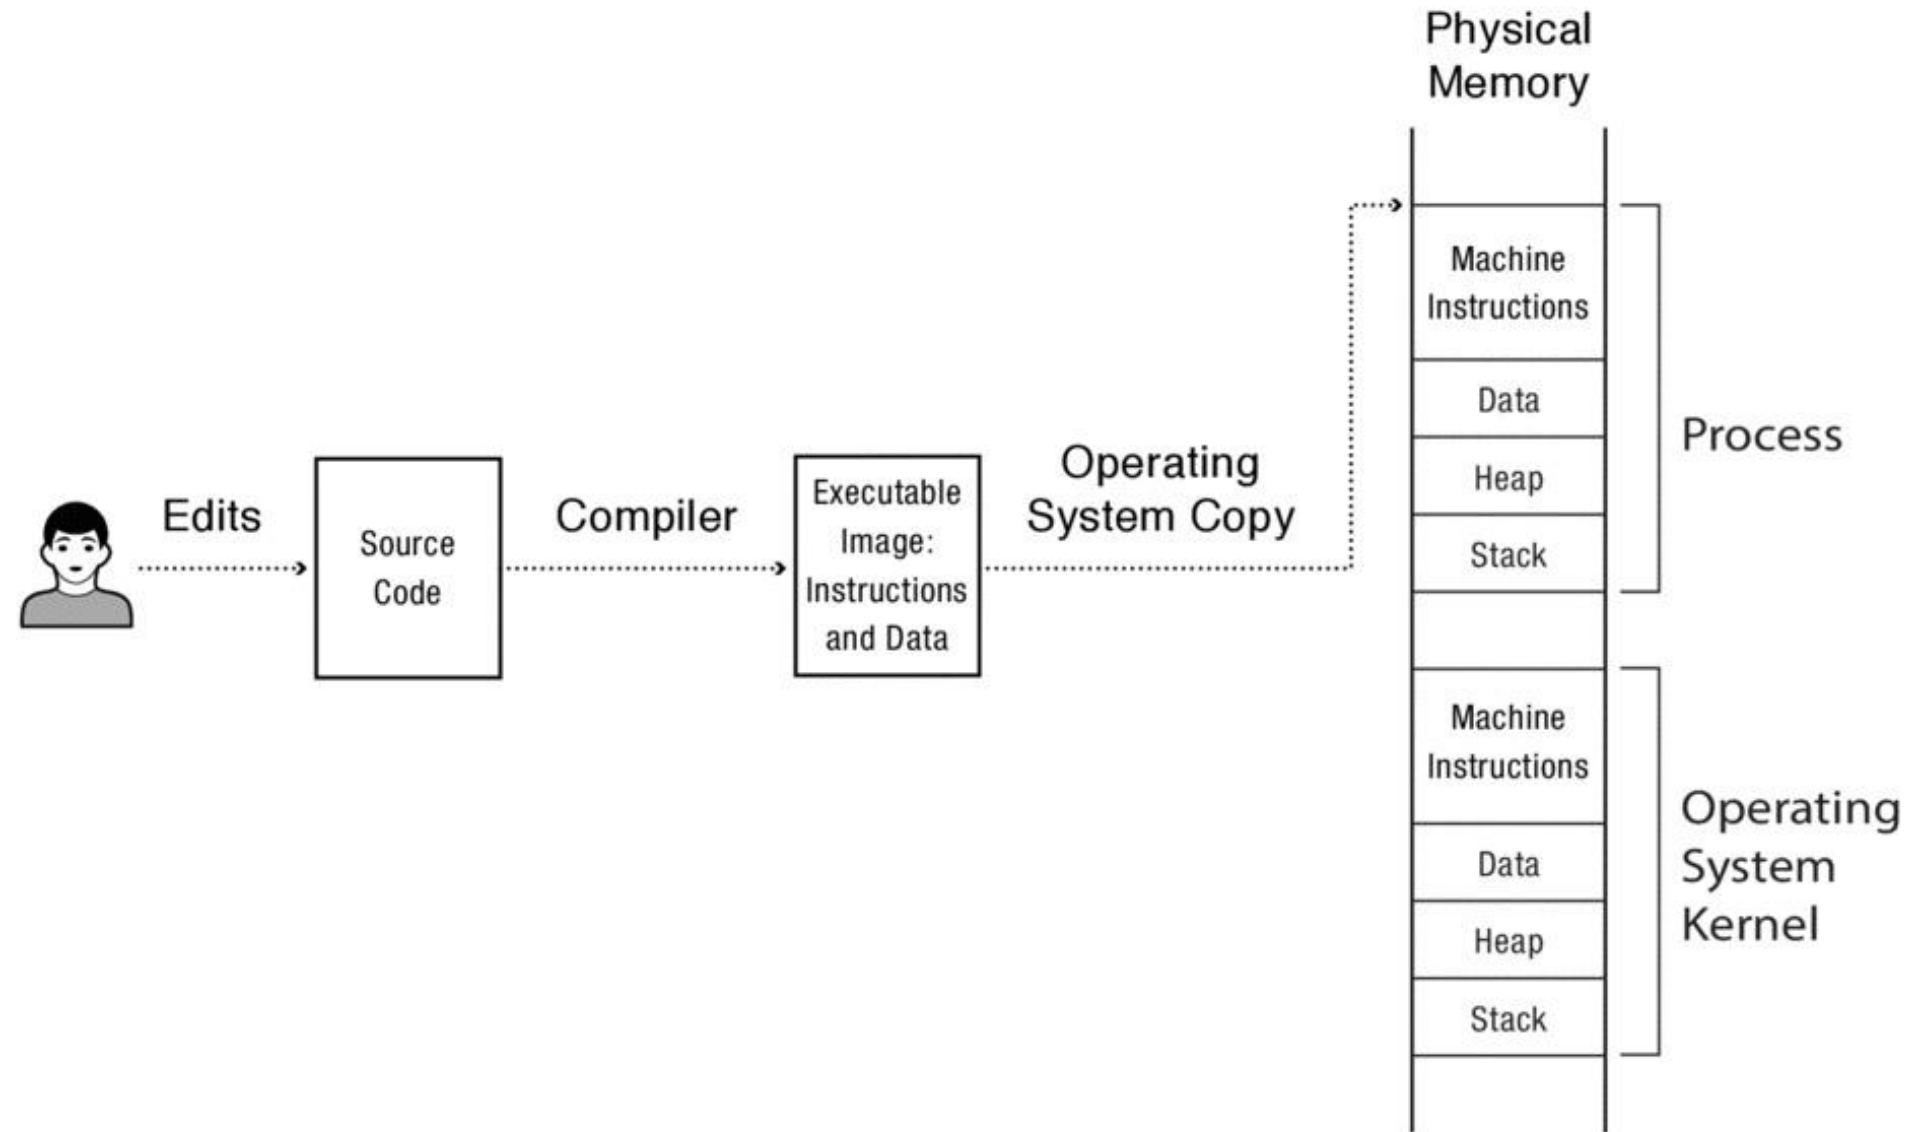
Program to Process

- We write a program in e.g. C
- A compiler turns that program into an instruction list
- The CPU interprets the instruction list (which is more a graph of basic blocks)

```
void add(int c) {  
    if(c == 1) {  
        ...  
    }  
    int main() {  
        int a = 2;  
        add(a);  
    }  
}
```

- Program becomes process when an executable file is loaded into memory
- Loading executable files via GUI mouse clicks, command line entry of its name

Program to Process

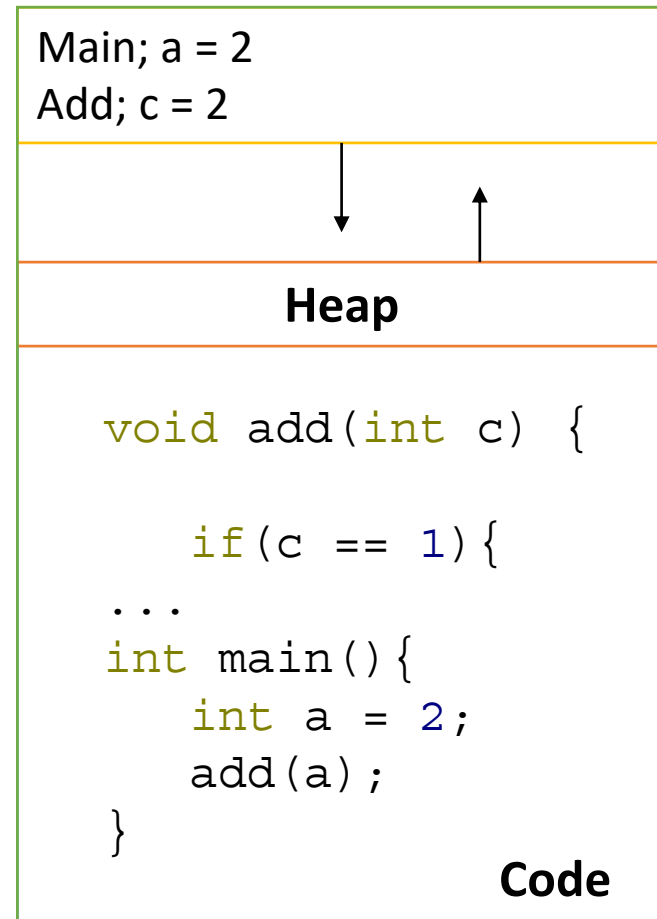


Process in Memory

- Program to process
 - What you wrote


```
void add(int c) {  
  
    if(c == 1) {  
    ...  
int main() {  
    int a = 2;  
    add(a);  
}
```

- What is in memory



Processes

- Process has its own address space:
 - The program code, also called **text section**
 - Current activity including program counter, processor registers
 - **Stack** containing temporary data
 - Function parameters
 - return addresses
 - local variables
 - **Data section** containing global variables
 - **Heap** containing memory dynamically allocated during run time

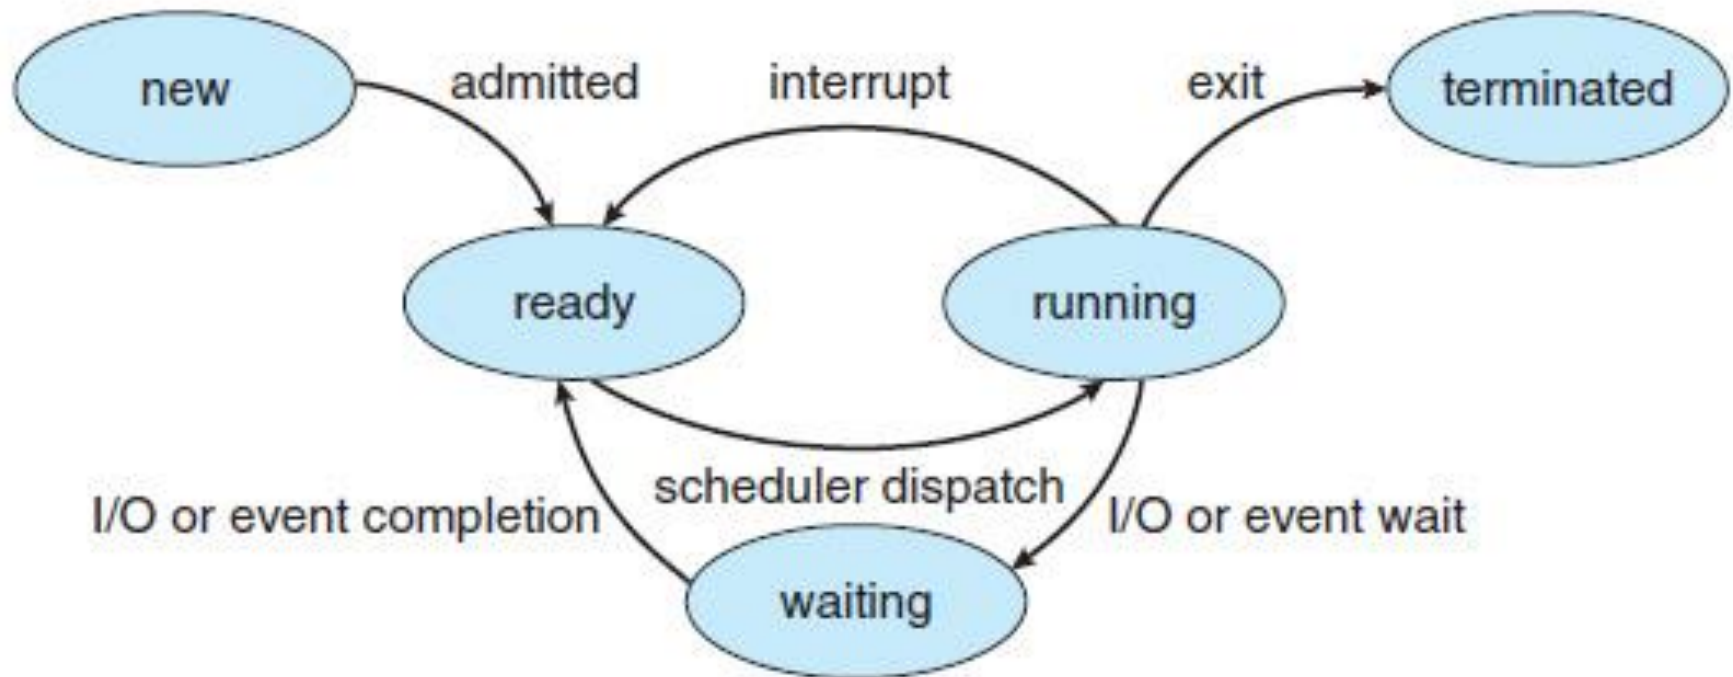


Process in Memory

Process State

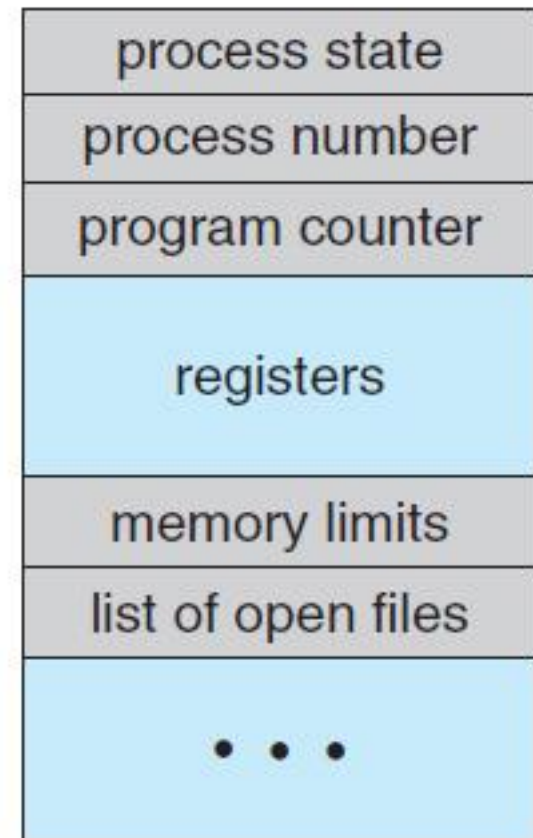
- As a process executes, it changes **state**
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a processor
 - **terminated**: the process has finished execution

Diagram of Process State

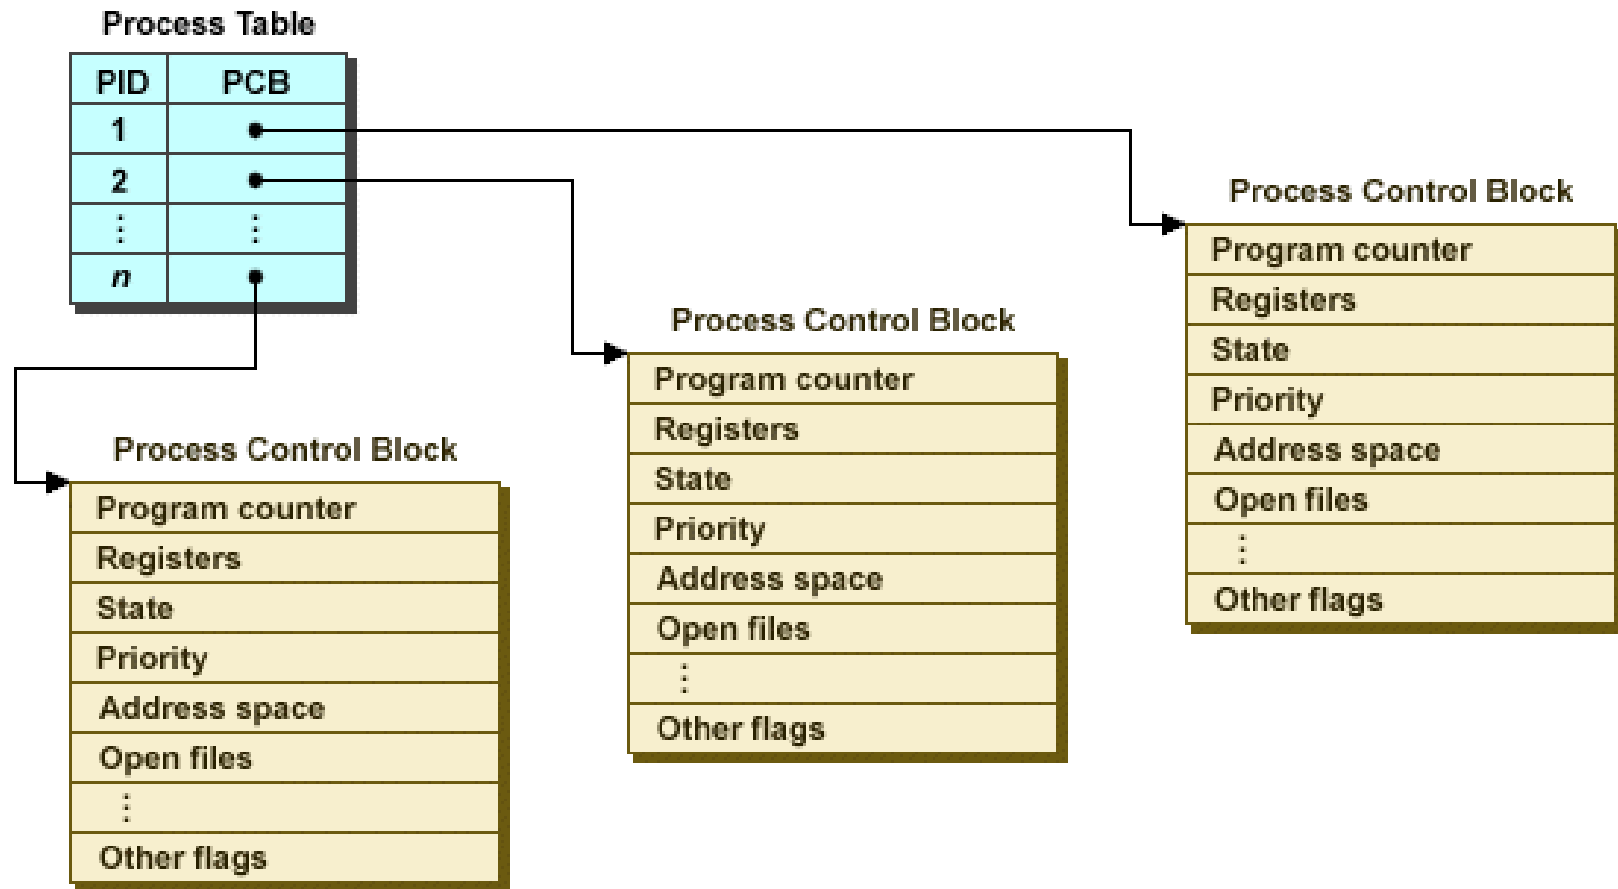


Process Control Blocks

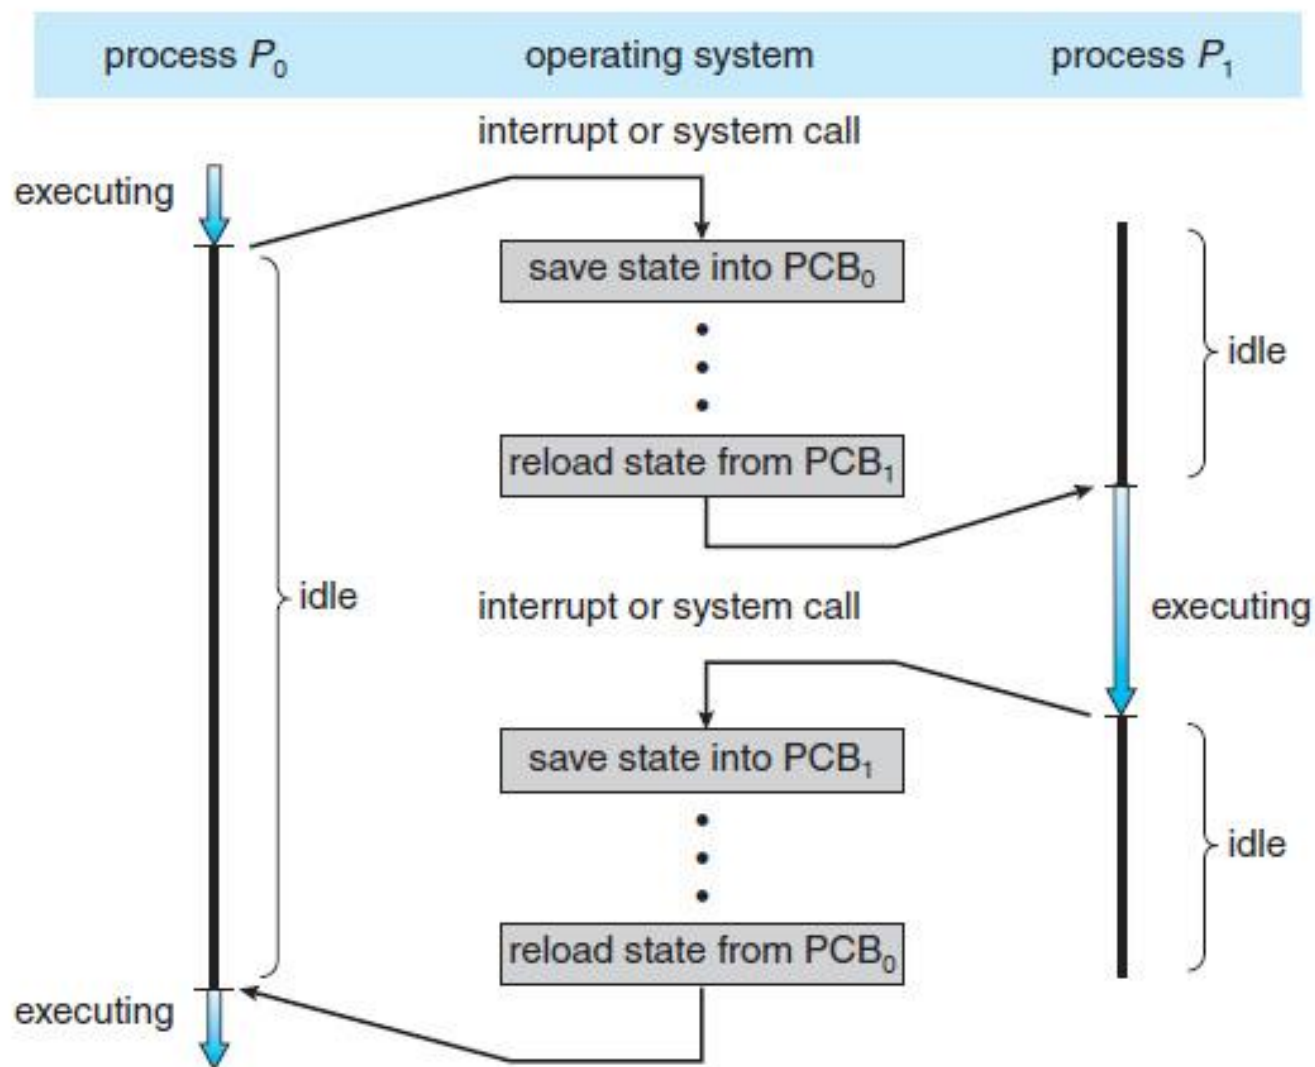
- Each process is represented in the operating system by a **process control block (PCB)**
- PCB contains many pieces of information associated with a specific process:
 - **Process state:** running, waiting, etc
 - **Program counter:** location of instruction to next execute
 - **CPU registers:** contents of all process-centric register
 - **CPU scheduling information:** priorities, scheduling queue pointers
 - **Memory-management information:** memory allocated to the process
 - **Accounting information:** CPU used, time limits, account numbers, job, or, process numbers
 - **I/O status information:** I/O devices allocated to process list of open files



Process Control Blocks



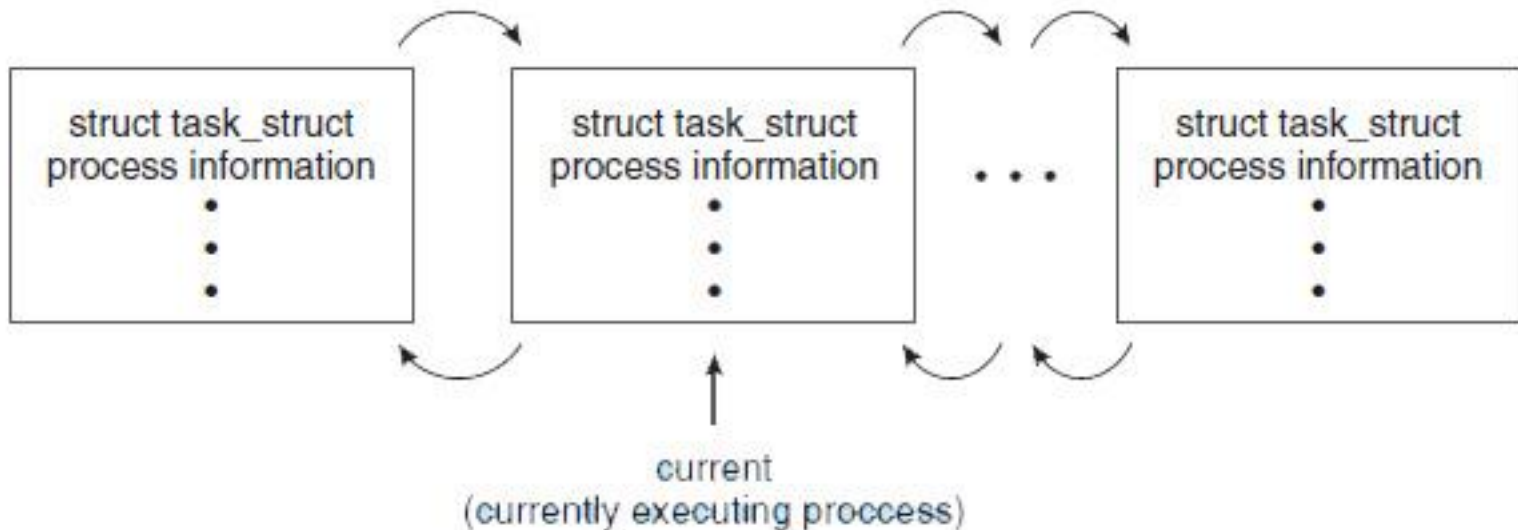
CPU Switch From Process to Process



Process Representation in Linux

- Represented by the C structure `task_struct`

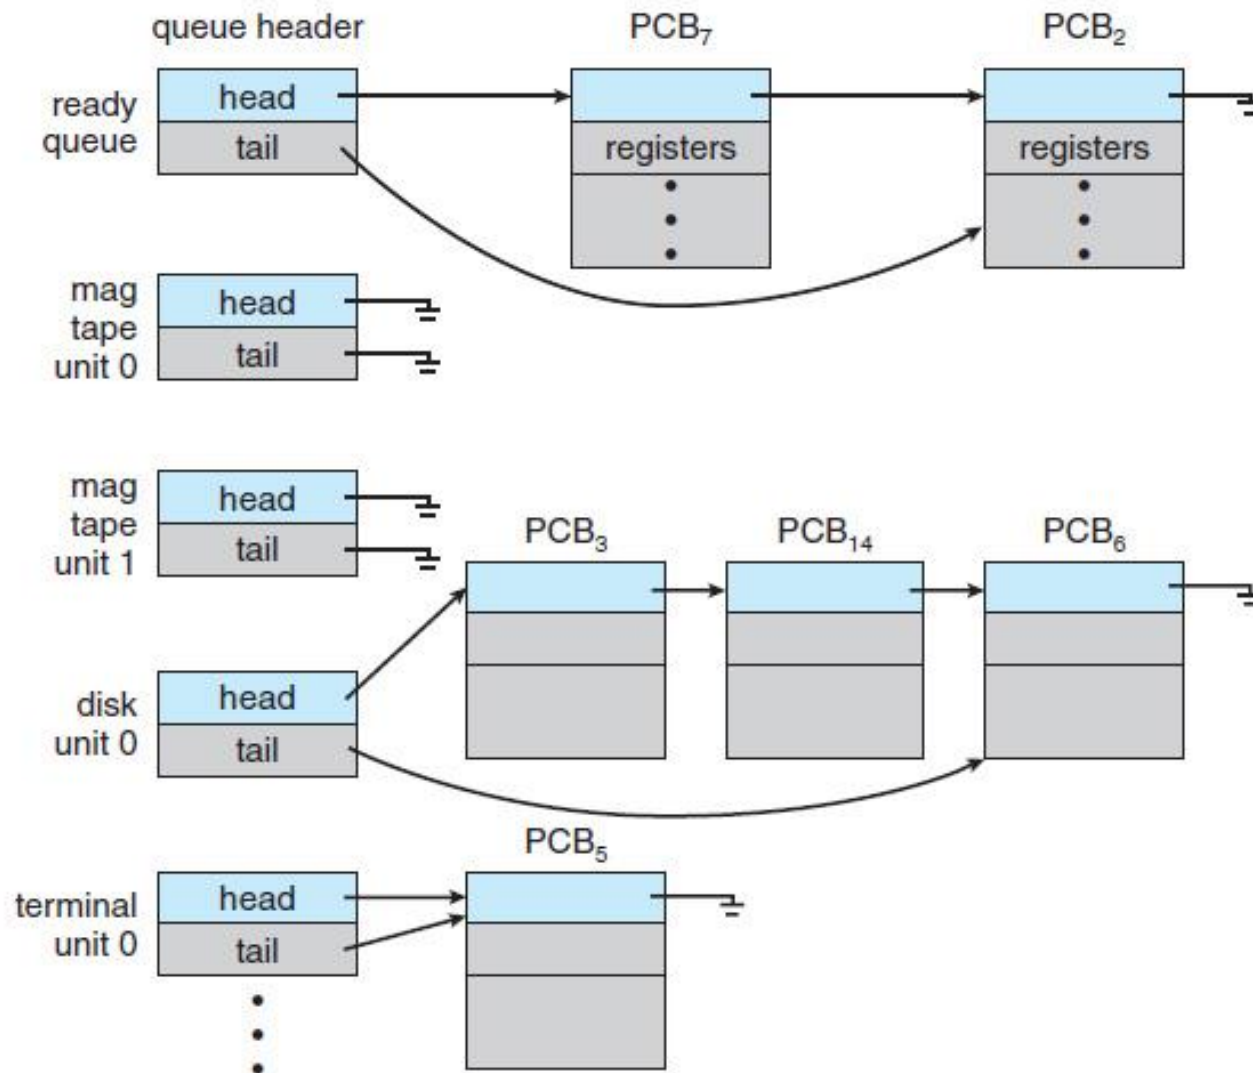
<code>pid_t</code> t_pid;	<code>/* process identifier */</code>
<code>long</code> state;	<code>/* state of the process */</code>
<code>unsigned int</code> time_slice	<code>/* scheduling information */</code>
<code>struct task_struct</code> *parent;	<code>/* this process's parent */</code>
<code>struct list_head</code> children;	<code>/* this process's children */</code>
<code>struct files_struct</code> *files;	<code>/* list of open files */</code>
<code>struct mm_struct</code> *mm;	<code>/* address space of this process */</code>



Process Scheduling

- Maximize CPU use, quickly switch processes onto CPU for time sharing
- **Process scheduler** selects among available processes for next execution on CPU
- Maintains scheduling queues of processes
 - **Job queue**: set of all processes in the system
 - **Ready queue**: set of all processes residing in main memory, read, and waiting to execute
 - **Device queues**: set of processes waiting for an I/O device
 - Processes migrate among the various queues

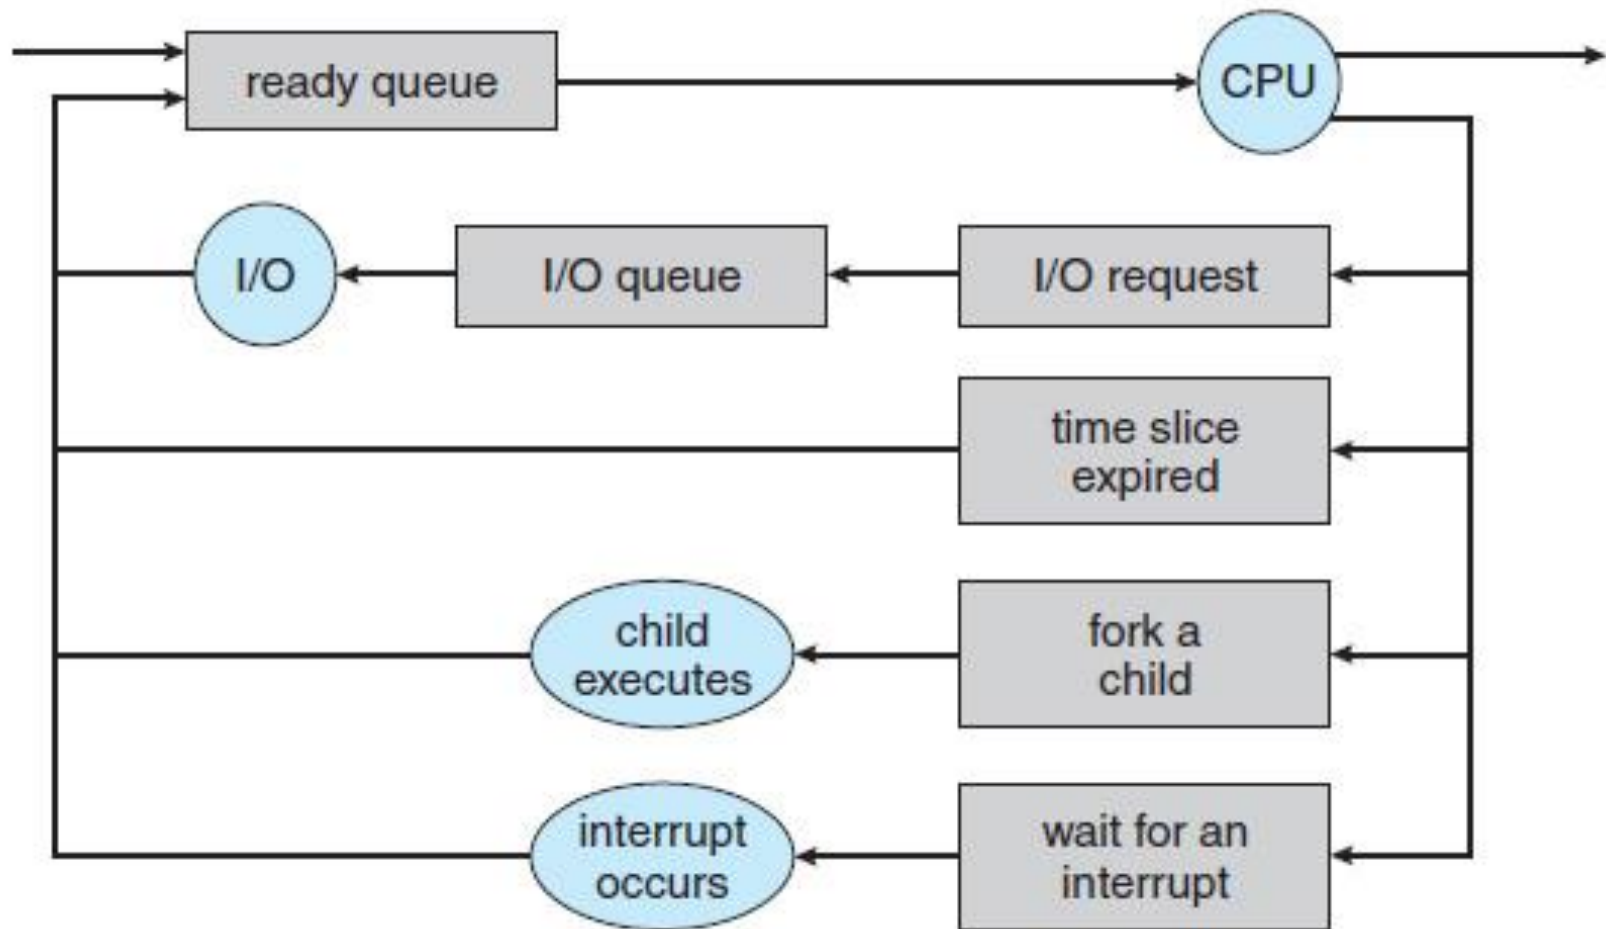
Ready Queue and Various I/O Device Queues



Scheduling Queues

- As processes enter the system, they are put into a **job queue** consisting of all processes in the system
- A new process is initially put in the ready queue
- It waits there until it is selected for execution
- Once the process is allocated the CPU and is executing, one of several events could occur:
 - The process could issue an I/O request and then be placed in an I/O queue
 - The process could create a new child process and wait for the child's termination
 - The process could be removed forcibly from the CPU, as a result of an interrupt, and be put back in the ready queue

Scheduling Queues



Schedulers

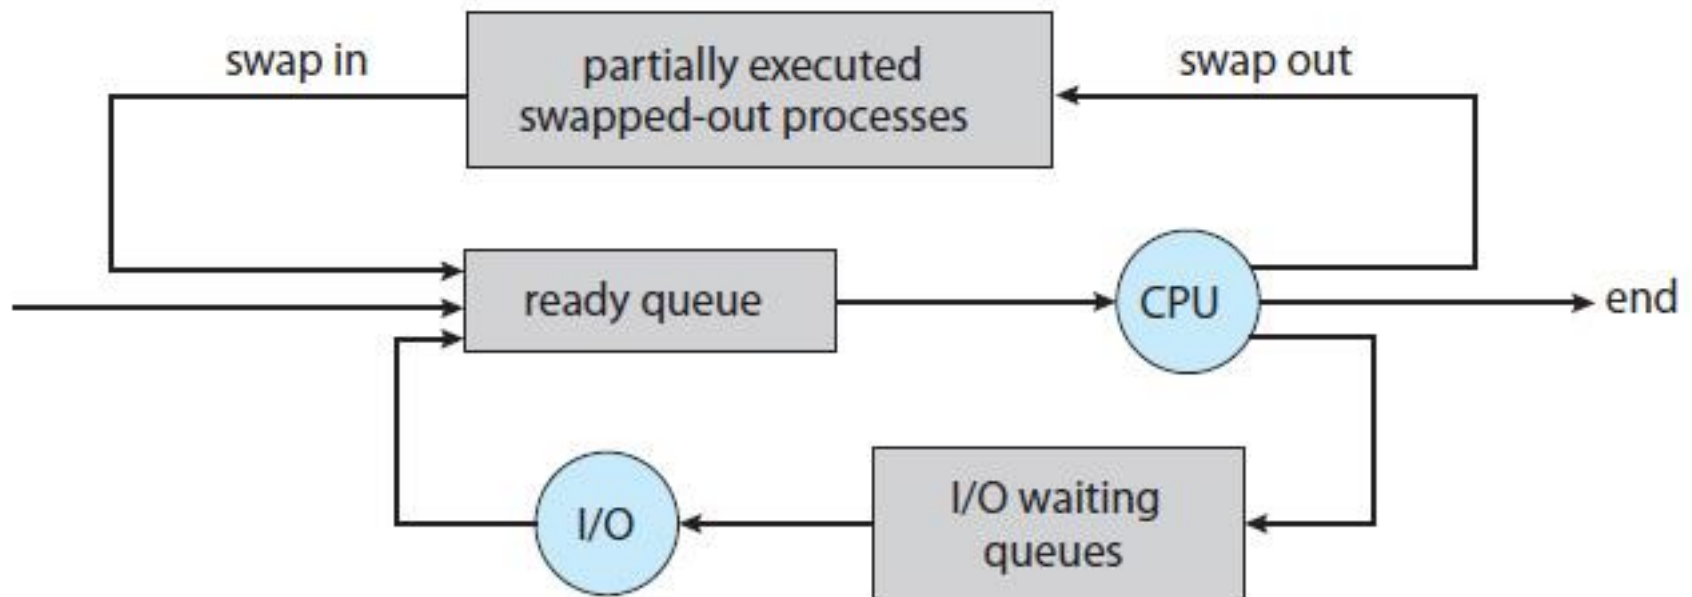
- A process migrates among the various scheduling queues throughout its lifetime
- The OS must select processes from the queues for scheduling
- The selection process is carried out by the appropriate **scheduler**
- There are three scheduler
 - Long-term scheduler or job scheduler
 - Short-term scheduler or CPU scheduler
 - Medium-term scheduler

Schedulers

- **Short-term scheduler** (or **CPU scheduler**): selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system
 - Short-term scheduler is invoked frequently (milliseconds) => (must be fast)
- **Long-term scheduler** (or **job scheduler**): selects which processes should be brought into the ready queue
 - Long-term scheduler is invoked infrequently (seconds, minutes) => (may be slow)
 - The long-term scheduler controls the **degree of multiprogramming**
- Processes can be described as either:
 - **I/O-bound process**: spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process**: spends more time doing computation; few very long CPU bursts
- Long-term scheduler strives for good ***process mix***

Medium Term Scheduling

- **Medium-term scheduler**: can be added if degree of multiple programming needs to decrease
 - Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**



Context Switch

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the PCB
- Context-switch time is overhead; the system does no useful work while switching
 - The more complex the OS and the PCB -> the longer the context switch
- Time dependent on hardware support
 - Some hardware provides multiple sets of registers per CPU

=> Multiple contexts loaded at once

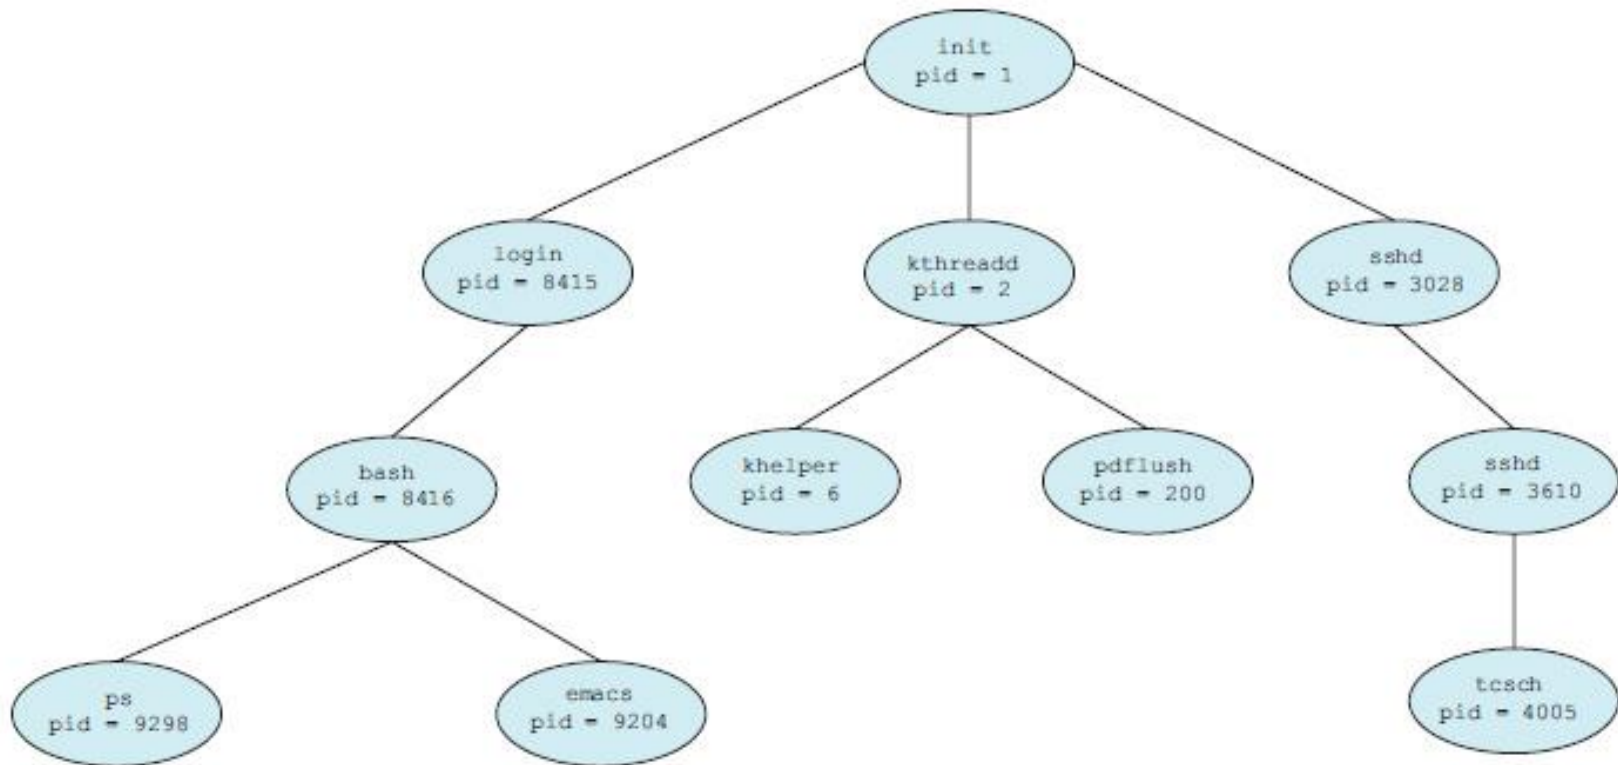
Operations in Processes

- Processes in most systems can execute concurrently
- They may be created and deleted dynamically
- System must provide mechanisms for:
 - Process creation
 - Process termination

Process Creation

- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via a **process identifier (pid)**
- Pid provides a unique value for each process in the system and can be used as an index to access various attributes of a process
- Resource sharing options:
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate

A Tree of Processes in Linux



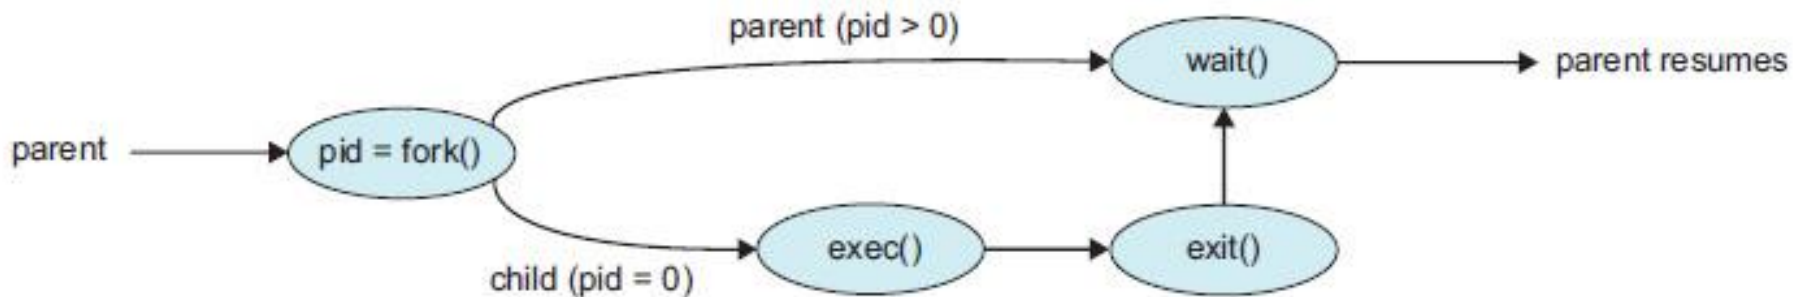
A Tree of Processes in Linux

- **Address space**

- Child duplication of parent
- Child has a program loaded into it

- Unix examples:

- **fork()** system call creates new process
- **exec()** system call used after a **fork()** to replace the process memory space with a new program



Creating a Separate Process using the Unix fork

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

Creating a Separate Process via Window API

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```

Process Termination

- Process executes last statement and then asks the operating system to delete it using the **exit()** system call
 - Returns status data from child to parent (via **wait()**)
 - Process resources are deallocated by operating system
- Parent may terminate the execution of children processes using the **abort()** system call
- Reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting and the operating systems does not allow a child to continue if its parent terminates

Process Termination

- Some operating systems do not allow child to exist if its parent has terminated
- If a process terminates, then all its children must also be terminated
 - **Cascading termination:** All children, grandchildren, etc. are terminated
 - The termination is initiated by the operating system
- The parent process may wait for termination of a child process by using the **wait()** system call
- The call returns status information and the pid of the terminated process

pid = wait(&status);

- If no parent waiting (did not invoke **wait()**) process is a **zombie**
- If parent terminated without invoking **wait**, process is an **orphan**